

Chapter 4

Phi-Regret and Multicalibration

Ioannis Anagnostides, Gabriele Farina, and Brian Hu Zhang*

The black-box reduction from Φ -regret to online learning due to [Gordon, Greenwald and Marks \[GGM08\]](#) centers around the algorithmic primitive of fixed points. This chapter introduces a different route to the same destination that passes through forecasting [\[FP26\]](#).

At a high level, the moral of this section is the following, also depicted in Figure 1.

One can construct a Φ -regret minimizer starting from a powerful enough *multicalibrated forecaster* of the upcoming utility, and best respond to such forecast. In turn, the task of constructing such a forecaster admits a reduction to online learning in the style of Gordon-Greenwald-Marks, with expected variational inequalities replacing expected fixed points as the nonlinear optimization primitive.

As we will show towards the end of the chapter, this forecasting-based reduction comes with algorithmic benefits that forego the need for complicated semiseparation (*cf.* Chapter 3, Section 3.2). We note however that it is not known whether a forecasting-based approach can yield fast *offline* algorithms for Φ -equilibrium computation (*cf.* Chapter 2, Section 2.3).




Figure 1: The two forecasting reductions and their relation to the Gordon-Greenwald-Marks framework.

*✉ ianagnos@cs.cmu.edu, {gfarina,zhangbh}@mit.edu. These tutorial notes have not undergone formal peer review. We are grateful for any feedback or reports of typos.

4.1 A Gordon-Greenwald-Marks result for multicalibration

We start with online multicalibration. In every round t , Nature reveals a context $\mathbf{c}^{(t)} \in \mathcal{C}$. The learner outputs a distribution $D^{(t)}$ over forecasts $\mathbf{p}^{(t)} \in \mathcal{U} \subseteq \mathbb{R}^d$, and Nature then reveals the true utility vector $\mathbf{u}^{(t)} \in \mathcal{U}$. Given test $h : \mathcal{C} \times \mathcal{U} \rightarrow \mathbb{R}^d$, the calibration error against h is

$$\text{MC-Err}^{(T)}(h) := \sum_{t=1}^T \mathbb{E}_{\mathbf{p}^{(t)} \sim D^{(t)}} [\langle h(\mathbf{c}^{(t)}, \mathbf{p}^{(t)}), \mathbf{u}^{(t)} - \mathbf{p}^{(t)} \rangle]. \quad (1)$$

The forecasts are $\mathcal{H}$ -multicalibrated if $\text{MC-Err}^{(T)}(h) = o(T)$ for every $h \in \mathcal{H}$.

As we show in Theorem 4.2, multicalibration admits a black-box reduction to online learning in the style of [Gordon, Greenwald and Marks \[GGM08\]](#). However, in the case of multicalibration the nonlinear primitive is *not* fixed points, but rather *expected variational inequalities*, as defined next.¹

Definition 4.1 (Expected variational inequality). Let $S : \mathcal{U} \rightarrow \mathbb{R}^d$ be an operator and let $\varepsilon > 0$. An ε -solution to the expected variational inequality induced by S is a distribution D over $\mathcal{U}$ such that

$$\mathbb{E}_{\mathbf{p} \sim D} [\langle S(\mathbf{p}), \mathbf{u} - \mathbf{p} \rangle] \leq \varepsilon \quad \forall \mathbf{u} \in \mathcal{U}. \quad (2)$$

Efficient algorithms for EVIs are known for general compact convex sets under mild oracle access assumptions [\[ZAT+25\]](#).

Armed with the EVI primitive, we can reduce multicalibration to an external-regret minimizer $R_{\mathcal{H}}$ whose decision set is the class of tests $\mathcal{H}$ as follows.

Algorithm: Multicalibration from external regret

Input: An external regret minimizer $R_{\mathcal{H}}$ for the test set $\mathcal{H}$

function NextForecast($\mathbf{c}^{(t)}$):

 Set $h^{(t)} := R_{\mathcal{H}}.\text{NextStrategy}()$

 Let $D^{(t)}$ solve the EVI with $S^{(t)}(\mathbf{p}) := h^{(t)}(\mathbf{c}^{(t)}, \mathbf{p})$:

$$\mathbb{E}_{\mathbf{p}^{(t)} \sim D^{(t)}} [\langle h^{(t)}(\mathbf{c}^{(t)}, \mathbf{p}^{(t)}), \mathbf{u} - \mathbf{p}^{(t)} \rangle] \leq \varepsilon^{(t)} \quad \forall \mathbf{u} \in \mathcal{U}$$

return $D^{(t)}$

function ObserveUtility($\mathbf{u}^{(t)}$):

 Feed to $R_{\mathcal{H}}$ the linear utility

$$g^{(t)}(h) := \mathbb{E}_{\mathbf{p}^{(t)} \sim D^{(t)}} [\langle h(\mathbf{c}^{(t)}, \mathbf{p}^{(t)}), \mathbf{u}^{(t)} - \mathbf{p}^{(t)} \rangle].$$

¹Expected variational inequalities have appeared in the literature under different names, including “outgoing minimax problems” [\[FH21\]](#), “accuracy certificates” [\[NOR10\]](#), or “negative correlation search” [\[PR25\]](#).

Theorem 4.2 ([FP26]). Let $\text{Reg}_{\mathcal{H}}^{(T)}(h)$ be the external regret of $R_{\mathcal{H}}$ with respect to the sequence of utilities $g^{(1)}, \dots, g^{(T)}$, and let $\text{EVI}^{(T)} := \sum_{t=1}^T \varepsilon^{(t)}$. Then the forecasts produced by the algorithm above satisfy

$$\text{MC-Err}^{(T)}(h) \leq \text{Reg}_{\mathcal{H}}^{(T)}(h) + \text{EVI}^{(T)} \quad \forall h \in \mathcal{H}.$$

Proof. The EVI condition is invoked at the realized utility vector $\mathbf{u}^{(t)}$, so

$$g^{(t)}(h^{(t)}) \leq \varepsilon^{(t)}.$$

Therefore, for any comparator test $h \in \mathcal{H}$,

$$\begin{aligned} \text{MC-Err}^{(T)}(h) &= \sum_{t=1}^T g^{(t)}(h) \\ &\leq \sum_{t=1}^T (g^{(t)}(h) - g^{(t)}(h^{(t)})) + \sum_{t=1}^T \varepsilon^{(t)} \\ &= \text{Reg}_{\mathcal{H}}^{(T)}(h) + \text{EVI}^{(T)}. \end{aligned}$$

Thus, sublinear regret over $\mathcal{H}$, together with $\sum^{(T)} \varepsilon^{(t)} = o(T)$, gives sublinear multicalibration error. $\square$

Necessity of EVIs. In a precise sense, EVIs are also *necessary*. Suppose we already had an efficient $\mathcal{H}$ -multicalibrated forecaster. Fix a context $\mathbf{c} \in \mathcal{C}$ and a test $h \in \mathcal{H}$, and consider the EVI operator $S(\mathbf{p}) = h(\mathbf{c}, \mathbf{p})$. To solve this EVI, run the forecaster for T rounds with the same context $\mathbf{c}^{(t)} = \mathbf{c}$. After it outputs $D^{(t)}$, choose Nature's response as

$$\mathbf{u}^{(t)} \in \arg \max_{\mathbf{u} \in \mathcal{U}} \left\langle \mathbf{u}, \mathbb{E}_{\mathbf{p} \sim D^{(t)}} [S(\mathbf{p})] \right\rangle. \quad (3)$$

For any fixed $\mathbf{u}^* \in \mathcal{U}$, the choice in (3) gives

$$\mathbb{E}_{\mathbf{p} \sim D^{(t)}} [\langle S(\mathbf{p}), \mathbf{u}^{(t)} - \mathbf{p} \rangle] \geq \mathbb{E}_{\mathbf{p} \sim D^{(t)}} [\langle S(\mathbf{p}), \mathbf{u}^* - \mathbf{p} \rangle].$$

Averaging the distributions $D^{(1)}, \dots, D^{(T)}$ uniformly therefore produces a distribution D with

$$\mathbb{E}_{\mathbf{p} \sim D} [\langle S(\mathbf{p}), \mathbf{u}^* - \mathbf{p} \rangle] \leq \frac{\text{MC-Err}^{(T)}(h)}{T} \quad \forall \mathbf{u}^* \in \mathcal{U}.$$

As T grows, the multicalibration condition drives the right-hand side to 0. So an online multicalibrated forecaster gives a black-box EVI solver.

4.2 From Multicalibration to Phi-regret minimization

We now use forecasting to construct a Φ -regret minimizer. Let $\mathcal{X} \subseteq \mathbb{R}^d$ be a compact convex action set, $\mathcal{U} \subseteq \mathbb{R}^d$ be a compact convex set of possible utility vectors, and $\Phi \subseteq \{\varphi : \mathcal{X} \rightarrow$

$\mathcal{X}$ be a family of deviations. Recall that if the learner outputs distributions $\mu^{(t)}$ over $\mathcal{X}$, its Φ -regret against a deviation $\varphi \in \Phi$ is

$$\Phi\text{Reg}^{(T)}(\varphi) := \sum_{t=1}^T \mathbb{E}_{\mathbf{x}^{(t)} \sim \mu^{(t)}} [\langle \varphi(\mathbf{x}^{(t)}) - \mathbf{x}^{(t)}, \mathbf{u}^{(t)} \rangle]. \quad (4)$$

The idea of the reduction asks the forecaster to predict the next utility vector. Given a forecast $\mathbf{p} \in \mathcal{U}$, define the deterministic best response

$$\sigma(\mathbf{p}) \in \arg \max_{\mathbf{x} \in \mathcal{X}} \langle \mathbf{x}, \mathbf{p} \rangle, \quad (5)$$

with ties broken by a fixed rule. If the forecaster outputs a distribution $D^{(t)}$ over forecasts $\mathbf{p}^{(t)}$, the decision maker plays the pushforward distribution $\mu^{(t)}$ induced by $\mathbf{x}^{(t)} = \sigma(\mathbf{p}^{(t)})$.

What tests should the forecaster be calibrated against? For each deviation $\varphi \in \Phi$, define

$$h_\varphi(\mathbf{p}) := \varphi(\sigma(\mathbf{p})) - \sigma(\mathbf{p}), \quad \mathcal{H}_\Phi := \{h_\varphi : \varphi \in \Phi\}. \quad (6)$$

As it turns out, multicalibration with respect to the test class $\mathcal{H}_\Phi$ suffices to control Φ -regret.

Algorithm: Φ -regret from multicalibration

Input: A forecaster over $\mathcal{U}$ multicalibrated with respect to $\mathcal{H}_\Phi$

NextStrategy():

 Query the forecaster and receive a distribution $D^{(t)}$ over forecasts $\mathbf{p}^{(t)} \in \mathcal{U}$
 return the pushforward distribution $\mu^{(t)}$ of $\mathbf{x}^{(t)} = \sigma(\mathbf{p}^{(t)})$

ObserveUtility($\mathbf{u}^{(t)}$):

 Feed the realized utility vector $\mathbf{u}^{(t)}$ to the forecaster

Theorem 4.3 ([FP26]). If the forecaster in the algorithm above has multicalibration error $\text{MC-Err}^{(T)}$ with respect to $\mathcal{H}_\Phi$, then the decision maker has

$$\Phi\text{Reg}^{(T)}(\varphi) \leq \text{MC-Err}^{(T)}(h_\varphi) \quad \forall \varphi \in \Phi.$$

In particular, sublinear $\mathcal{H}_\Phi$ -multicalibration implies sublinear Φ -regret.

Proof. Fix any $\varphi \in \Phi$. Since $\mu^{(t)}$ is the pushforward of $D^{(t)}$ under $\mathbf{p} \mapsto \sigma(\mathbf{p})$,

$$\begin{aligned} \Phi\text{Reg}^{(T)}(\varphi) &= \sum_{t=1}^T \mathbb{E}_{\mathbf{p}^{(t)} \sim D^{(t)}} [\langle h_\varphi(\mathbf{p}^{(t)}), \mathbf{u}^{(t)} \rangle] \\ &= \sum_{t=1}^T \mathbb{E}_{\mathbf{p}^{(t)} \sim D^{(t)}} [\langle h_\varphi(\mathbf{p}^{(t)}), \mathbf{u}^{(t)} - \mathbf{p}^{(t)} \rangle] + \sum_{t=1}^T \mathbb{E}_{\mathbf{p}^{(t)} \sim D^{(t)}} [\langle h_\varphi(\mathbf{p}^{(t)}), \mathbf{p}^{(t)} \rangle]. \end{aligned}$$

The first term is exactly $\text{MC-Err}^{(T)}(h_\varphi)$. The second term is nonpositive, because

$$\langle h_\varphi(\mathbf{p}), \mathbf{p} \rangle = \langle \varphi(\sigma(\mathbf{p})), \mathbf{p} \rangle - \langle \sigma(\mathbf{p}), \mathbf{p} \rangle \leq 0$$

by the definition of $\sigma(\mathbf{p})$ as a maximizer over $\mathcal{X}$. □

Quite intuitively, the complexity of the required calibration class scales with the complexity of the deviation class. We can think of Theorem 4.3 as a generalization of the classical connection between ℓ_1 -calibration and swap regret [FV97].

Remark 4.4. We can equally easily define a contextual version of the result. If contexts $\mathbf{c}^{(t)}$ are observed before play and deviations have the form $\varphi : \mathcal{C} \times \mathcal{X} \rightarrow \mathcal{X}$, then we define

$$h_\varphi(\mathbf{c}, \mathbf{p}) := \varphi(\mathbf{c}, \sigma(\mathbf{p})) - \sigma(\mathbf{p}).$$

Multicalibration with respect to these tests gives a notion of *contextual* Φ -regret.

4.2.1 Putting the two reductions together

Combining Theorem 4.2 and Theorem 4.3 gives the promised route from external regret to Φ -regret. Set $\mathcal{H} = \mathcal{H}_\Phi$. Run the multicalibration algorithm as the forecaster inside the best-response algorithm. Then, for every $\varphi \in \Phi$,

$$\Phi\text{Reg}^{(T)}(\varphi) \leq \text{MC-Err}^{(T)}(h_\varphi) \leq \text{Reg}_{\mathcal{H}}^{(T)}(h_\varphi) + \text{EVI}^{(T)}. \quad (7)$$

Thus, the burden of Φ -regret minimization shifts to two primitives:

1. no-regret learning over the induced test class $\mathcal{H}_\Phi$; and
2. solving EVIs over the forecast domain $\mathcal{U}$.

The key distinction from the Gordon-Greenwald-Marks path is that we do *not* need to optimize over valid deviations directly. We only need a regret minimizer over tests that contain the maps h_φ . This extra flexibility is useful when characterizing the deviation class would otherwise require complex machinery, such the semiseparation construction of Daskalakis, Farina, Fishelson, Pipis and Schneider [DFF+25] mentioned in Chapter 2, Section 2.2.

4.2.2 Linear deviations

As an illustration, consider linear deviations $\varphi(\mathbf{x}) = \mathbf{M}\mathbf{x}$ over a convex compact set $\mathcal{X} \subseteq \mathbb{R}^d$. The classical GGM approach asks us to understand the geometry of all matrices $\mathbf{M}$ satisfying $\mathbf{M}\mathcal{X} \subseteq \mathcal{X}$, and then compute fixed points of the matrices selected by the regret minimizer. That geometry of the endomorphisms can be difficult [DFF+25].

The multicalibration route permits a relaxation. From (6),

$$h_\varphi(\mathbf{p}) = (\mathbf{M} - \mathbf{I})\sigma(\mathbf{p}).$$

If every valid linear endomorphism has spectral norm at most S , then each test above lies in a Frobenius ball of radius on the order of $\sqrt{d}(S + 1)$. Hence, instead of learning over valid endomorphisms, we can learn over the larger class

$$\mathcal{H} := \{\mathbf{p} \mapsto \mathbf{A}\sigma(\mathbf{p}) : \|\mathbf{A}\|_F \leq \rho\},$$

where ρ is chosen large enough to contain all tests h_φ . This larger class is easy: external regret over a Euclidean ball can be easily handled by projected online gradient descent. Indeed, recall that the utility sent to the matrix learner at time t is linear:

$$g^{(t)}(\mathbf{A}) = \mathbb{E}_{\mathbf{p}^{(t)} \sim D^{(t)}} [\langle \mathbf{A} \sigma(\mathbf{p}^{(t)}), \mathbf{u}^{(t)} - \mathbf{p}^{(t)} \rangle] = \left\langle \mathbf{A}, \mathbb{E}_{\mathbf{p}^{(t)} \sim D^{(t)}} \left[(\mathbf{u}^{(t)} - \mathbf{p}^{(t)}) \sigma(\mathbf{p}^{(t)})^\top \right] \right\rangle_F.$$

Consequently, if $\|\mathbf{x}\|_2 \leq B$ for $\mathbf{x} \in \mathcal{X}$ and $\|\mathbf{u}\|_2 \leq L$ for $\mathbf{u} \in \mathcal{U}$, the standard projected-gradient bound gives $O(BLS\sqrt{dT})$ linear-swap regret, up to constants. This improves the dimension dependence of the recent semi-separation-based route of [Daskalakis, Farina, Fishelson, Pipis and Schneider \[DFF+25\]](#), while avoiding semi-separation altogether.

4.2.3 RKHS deviations

The same idea extends beyond finite-dimensional linear classes. Suppose that the deviation functions $\varphi : \mathcal{C} \times \mathcal{X} \rightarrow \mathcal{X}$ lie in a vector-valued reproducing kernel Hilbert space with matrix-valued kernel

$$\Gamma((\mathbf{c}, \mathbf{x}), (\mathbf{c}', \mathbf{x}')) \in \mathbb{R}^{d \times d}.$$

For the contextual reduction, the relevant tests are

$$h_\varphi(\mathbf{c}, \mathbf{p}) = \varphi(\mathbf{c}, \sigma(\mathbf{p})) - \sigma(\mathbf{p}).$$

These tests also lie in an RKHS. Indeed, the first term is represented by composing the deviation kernel with the best-response map, while the second term is represented by the linear kernel on $\sigma(\mathbf{p})$. The resulting kernel is

$$\Gamma'((\mathbf{c}, \mathbf{p}), (\mathbf{c}', \mathbf{p}')) := \Gamma((\mathbf{c}, \sigma(\mathbf{p})), (\mathbf{c}', \sigma(\mathbf{p}'))) + \sigma(\mathbf{p})\sigma(\mathbf{p}')^\top. \quad (8)$$

So, if we have an online learner over the RKHS ball associated with Γ' , Theorem 4.2 supplies a multicalibrated forecaster, and Theorem 4.3 turns it into a no- Φ -regret algorithm. This recovers low-degree polynomial deviations as a special case and also covers infinite-dimensional kernels, such as Gaussian kernels, where the fixed-point route does not have an obvious finite-dimensional endomorphism geometry to exploit.

Bibliography for this chapter

- [GGM08] G. J. Gordon, A. Greenwald and C. Marks. (2008). No-regret learning in convex games. *International Conference on Machine Learning*.
- [FP26] G. Farina and J. C. Perdomo. (2026). An Efficient Black-Box Reduction from Online Learning to Multicalibration, and a New Route to Φ -Regret Minimization. *arXiv preprint arXiv:2604.19592*.
- [FH21] D. P. Foster and S. Hart. (2021). Forecast Hedging and Calibration. *Journal of Political Economy*.

- [NOR10] A. Nemirovski, S. Onn and U. G. Rothblum. (2010). Accuracy Certificates for Computational Problems with Convex Structure. *Mathematics of Operations Research*.
- [PR25] J. C. Perdomo and B. Recht. (2025). In Defense of Defensive Forecasting. *arXiv preprint arXiv:2506.11848*.
- [ZAT+25] B. H. Zhang, I. Anagnostides, E. Tewolde, R. E. Berker, G. Farina, V. Conitzer and T. Sandholm. (2025). Expected Variational Inequalities. *International Conference on Machine Learning*, 74422–74446.
- [FV97] D. P. Foster and R. V. Vohra. (1997). Calibrated learning and correlated equilibrium. *Games and Economic Behavior*, 21, 40–55.
- [DFF+25] C. Daskalakis, G. Farina, M. Fishelson, C. Pipis and J. Schneider. (2025). Efficient Learning and Computation of Linear Correlated Equilibrium in General Convex Games. *Symposium on Theory of Computing (STOC)*.