

Chapter 5

TreeSwap: Efficient Swap Regret Minimization

Ioannis Anagnostides, Gabriele Farina, and Brian Hu Zhang*

my intention so far is for this chapter to come last, after even the multicalibration section -brian

So far, we have seen that no-regret learning, and efficient equilibrium computation, are essentially possible whenever the set of deviations Φ is “efficiently representable”, in particular, whenever *external* regret minimization is possible over Φ . What about beyond? In particular, we have not yet addressed the *strongest* possible notion of Φ -regret, namely, the case where Φ contains *all* functions $\varphi : \mathcal{X} \rightarrow \mathcal{X}$ —known as *swap regret*.

The fact that algorithms for Φ -regret minimization grow increasingly complex as Φ itself grows in complexity may lead one to believe that efficient swap regret minimization is hopeless. For example, the set of all functions $\varphi : \mathcal{X} \rightarrow \mathcal{X}$ has *knfinite* dimension, so the techniques we have covered so far do not work. If $\mathcal{X}$ is a polytope, we could run the Blum-Mansour algorithm over its vertices—this would indeed succeed, but it has the drawback that the number of vertices of $\mathcal{X}$ is generally not polynomial in the dimension d of $\mathcal{X}$, and hence the algorithm is not efficient. Hence, for a long time, the possibility of swap regret minimization—and correspondingly, correlated equilibrium computation beyond normal-form games—remained a major open question.

The question has since essentially been fully resolved. In a simultaneous major breakthrough, Peng and Rubinstein [PR24] there's still a extra period before the []. also, is it possible to have first names/initials in the references section but not in inline citations? and Dagan, Daskalakis, Fishelson and Golowich [Dag+24], using essentially an identical algorithm now known as TreeSwap, showed that, if there is a regret minimizer on $\mathcal{X}$ that achieves *external* regret ϵ after M rounds, then there is a regret minimizer on $\mathcal{X}$ that achieves ϵT *swap* regret after $M^{1/\epsilon}$ rounds, and therefore efficient swap regret minimization *ks* in fact possible, at least when ϵ is a constant. They also showed nearly-matching lower bounds for normal-form games, which Daskalakis, Farina, Golowich, Sandholm and Zhang [Das+24] later extended to extensive-form games as well, thereby precluding the possibility of $\text{poly}(d, 1/\epsilon)$ -time algorithms for swap regret beyond normal-form games. In this chapter, we will go over the upper bound.

*✉ ianagnos@cs.cmu.edu, {gfarina,zhangbh}@mit.edu. These tutorial notes have not undergone formal peer review. We are grateful for any feedback or reports of typos.

5.1 The TreeSwap Algorithm

We now describe the TreeSwap algorithm, following Peng and Rubinstein [PR24] and Dagan, Daskalakis, Fishelson and Golowich [Dag+24]. Let $\mathcal{R}$ be a no-regret algorithm on $\mathcal{X}$ whose regret is bounded by ϵT after M rounds. Of course, $\mathcal{R}$ will not necessarily have small swap regret. So, how can we hope to bound swap regret?

One thing to try might be the following. Instead of having $\mathcal{R}$ update its strategy on *every* iteration, it will be lazy, and only update its strategy every M iterations, using the average utility on those M iterations. That is, at time M , it receives utility $\frac{1}{M} \sum_{t=1}^M \mathbf{u}^{(t)}$, at time $2M$ it receives utility $\frac{1}{M} \sum_{t=M+1}^{2M} \mathbf{u}^{(t)}$, and so on. Then, between those M iterations, that is, while $\mathcal{R}$ is playing the same strategy, we initialize *another copy of the same regret minimizer*, which we will call $\mathcal{R}'$, which updates its strategy on *every* iteration.

The purpose of $\mathcal{R}'$ is to bound the value of the most profitable deviation from every single strategy that $\mathcal{R}$ plays. Therefore, $\mathcal{R}'$ can ensure that no strategy played by $\mathcal{R}$ incurs large swap regret. Of course, $\mathcal{R}'$ itself will have high swap regret... but this can be solved by introducing yet more layers!

The TreeSwap algorithm uses K layers of regret minimizers $\mathcal{R}_0, \dots, \mathcal{R}_{K-1}$, each responsible for ensuring that there is no profitable swap deviation for the regret minimizers in the next layer. $\mathcal{R}_0$ updates on every iteration. Each regret minimizer $\mathcal{R}_k$ updates M times slower than the previous regret minimizer $\mathcal{R}_{k-1}$, and resets every M updates. Hence, this algorithm runs for M^K steps. The strategy output by TreeSwap is the uniform distribution (not to be confused with the average) over the strategies $(\mathbf{x}_0^{(t)}, \dots, \mathbf{x}_{K-1}^{(t)})$ output by the K regret minimizers. For simplicity of notation, we will henceforth zero-index timesteps and use the notation $[n] = \{0, \dots, n-1\}$.

Algorithm: TreeSwap

Input: Regret minimizers $\mathcal{R}_1, \dots, \mathcal{R}_K$

for timesteps $t \in [T] := [M^K]$ **do**

for layers $k \in [K]$ **do**

if $M^{k+1} \mid t$ **then** reset $\mathcal{R}_k$

else if $M^k \mid t$ **then** pass to $\mathcal{R}_k$ the average of the last M^k utilities,

$$\frac{1}{M^k} \sum_{\tau=t-M^k}^{t-1} \mathbf{u}^{(\tau)}$$

$\mathbf{x}_k^{(t)} \leftarrow$ current strategy of $\mathcal{R}_k$

 play mixed strategy $\text{Unif}(\mathbf{x}_0^{(t)}, \dots, \mathbf{x}_{K-1}^{(t)}) \in \Delta(\mathcal{X})$

 receive utility $\mathbf{u}^{(t)} \in \mathbb{R}^d$

Theorem 5.1. The regret of TreeSwap is bounded by $T \cdot (\epsilon + 1/K)$. In particular, taking $K = 1/\epsilon$ and $M = \text{poly}(d, 1/\epsilon)$ (as is achieved by any reasonable external-regret-minimizing algorithm), TreeSwap achieves swap regret ϵT after $M^K = d^{\tilde{O}(1/\epsilon)}$ iterations.

Remark 5.2. Theorem 5.1 only applies when T is a power of M . Handling the case where T is not a power of M requires a bit of care and results in a looser bound of $T \cdot (\epsilon + 3/K)$, but the main ideas are captured by the above result and its proof. [Dag+24]

can [Dag+24] be displayed as [DDFS24] instead? I think `\alpha`` style in LaTeX does [ABC+00] for ≥ 4 authors

Remark 5.3. We are mainly interested in the implications of Theorem 5.1 for its implications for swap regret in high-dimensional settings. However, it is worth noting that Theorem 5.1 achieves a new result even in the normal-form setting: its regret bound for normal form, when instantiated with multiplicative weights, is $\log(N)^{\tilde{O}(1/\epsilon)}$, which, for large N and ϵ , is better than the bound of Blum-Mansour.

We now sketch a proof of Theorem 5.1.

Proof Sketch. In this proof, timesteps are zero-indexed, and $[n] := \{0, \dots, n-1\}$. Let $\varphi : \mathcal{X} \rightarrow \mathcal{X}$ be any function. Then the (time-averaged) regret against φ is given by

$$\begin{aligned} \frac{1}{T} \text{Reg}(T, \varphi) &= \frac{1}{K} \sum_{k=0}^{K-1} \mathbb{E}_{t \in [T]} \left\langle \mathbf{u}^{(t)}, \varphi(\mathbf{x}_k^{(t)}) \right\rangle - \left\langle \mathbf{u}^{(t)}, \mathbf{x}_k^{(t)} \right\rangle \\ &= \frac{1}{K} \sum_{k=0}^{K-1} \mathbb{E}_{t \in [T]} \left\langle \mathbf{u}^{(t)}, \varphi(\mathbf{x}_k^{(t)}) \right\rangle - \frac{1}{K} \sum_{k=1}^K \mathbb{E}_{t \in [T]} \left\langle \mathbf{u}^{(t)}, \mathbf{x}_{k-1}^{(t)} \right\rangle \\ &= \frac{1}{K} \sum_{k=1}^{K-1} \underbrace{\mathbb{E}_{t \in [T]} \left\langle \mathbf{u}^{(t)}, \varphi(\mathbf{x}_k^{(t)}) - \mathbf{x}_{k-1}^{(t)} \right\rangle}_{\leq \epsilon} + \frac{1}{K} \underbrace{\mathbb{E}_{t \in [T]} \left\langle \mathbf{u}^{(t)}, \varphi(\mathbf{x}_0^{(t)}) - \mathbf{x}_{K-1}^{(t)} \right\rangle}_{\leq 1} \\ &\leq \epsilon + \frac{1}{K} \end{aligned}$$

where, in the second-to-last line, the bound on the first term comes from the regret bound of each phase between resets (of length M updates), noting that $\varphi(\mathbf{x}_k^{(t)})$ is constant within any given phase and thus the regret against $\varphi(\mathbf{x}_k^{(t)})$ is bounded by the external regret during the phase. $\square$

Bibliography for this chapter

- [PR24] B. Peng and A. Rubinstein. (2024). Fast Swap Regret Minimization and Applications to Approximate Correlated Equilibria. *Symposium on Theory of Computing (STOC)*.
- [Dag+24] Y. Dagan, C. Daskalakis, M. Fishelson and N. Golowich. (2024). From External to Swap Regret 2.0: An Efficient Reduction for Large Action Spaces. *Symposium on Theory of Computing (STOC)*.

- [Das+24] C. Daskalakis, G. Farina, N. Golowich, T. Sandholm and B. H. Zhang. (2024).
A Lower Bound on Swap Regret in Extensive-Form Games. *arXiv preprint
arXiv:2406.13116*.